

Reinforcement Learning 02


DS Development Group Presentations

Peter Nicholas S. Onglao | MNL

Background

GOAL:

Make an RL algorithm for chess (>2 presentations away) that can beat me



- Planning by Dynamic Programming (DP)
- Monte Carlo (MC) Learning
- Temporal Difference (TD) Learning

- [UCL Course on Reinforcement Learning](#), David Silver
- [Reinforcement Learning – An Introduction](#), Sutton and Barto
- [Reinforcement Learning \(GitHub repo\)](#), Denny Britz

1

Dynamic Programming

When the MDP is known



Dynamic Programming

- Dynamic – sequential/temporal component
- Programming – optimizing a “program” (policy)
- Break down complex problems into subproblems
 - ▷ Solve the subproblems
 - ▷ Combine solutions to subproblems



Dynamic Programming

- DP is a general solution method for problems with two properties:
 - ▷ Optimal substructure
 - ▷ *Principle of optimality* applies
 - ▷ Optimal solution can be decomposed into subproblems
 - ▷ Bellman equation – recursive decomposition
 - ▷ Overlapping subproblems
 - ▷ Subproblems recur many times
 - ▷ Solutions can be cached and reused
 - ▷ Value function



Dynamic Programming

- Assumes **full knowledge** of the MDP
- Used for *planning*
- Prediction
 - Input: MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ and **policy π**
 - Output: **value function v_π**
- Control
 - Input: MDP $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
 - Output: **optimal value function v_*** and **optimal policy π_***



Generalized Policy Iteration (GPI)

■ Given a policy π :

▷ Evaluate the policy π

$$v_{\pi}(s) = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \dots | S_t = s]$$

▷ Improve the policy by acting greedily w/r v_{π}

$$\pi' = \text{greedy}(v_{\pi})$$


■ GPI: let policy-evaluation and policy-improvement interact

■ Most RL methods are well described as GPI

■ This process of policy iteration always converges to π^*



Generalized Policy Iteration (GPI)




Policy evaluation: Estimate v_π

Iterative policy evaluation

Policy improvement: Generate $\pi' \geq \pi$

Greedy policy improvement





Principle of Optimality

- An optimal policy can be subdivided into two components:
 - ▷ An optimal first action A_*
 - ▷ Followed by an optimal policy from successor state S'
- Theorem (Principle of Optimality)
*A policy $\pi(a|s)$
achieves the optimal value from state s , $v_\pi(s) = v_*(s)$,
iff for any state s' reachable from s ,
 π achieves the optimal value from state s' , $v_\pi(s') = v_*(s')$*



Deterministic Value Iteration

- If we know the solution to subproblems $v_*(s')$
- Then solution $v_*(s)$ can be found by **one-step lookahead**

$$v_*(s) \leftarrow \max_{a \in \mathcal{A}} \left(\mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_*(s') \right)$$

- Value iteration: apply these updates iteratively
- Start with final rewards and work backwards



Synchronous DP Algorithms

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Iterative Policy Evaluation
Control	Bellman Expectation Equation + Greedy Policy Improvement	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

- Algorithms are based on state-value function $v_{pi}(s)$ or $v_*(s)$
- Complexity $O(mn^2)$ per iteration, for m actions and n states
- Could also apply to action-value function $q_\pi(s, a)$ or $q_*(s, a)$
- Complexity $O(m^2n^2)$ per iteration



Asynchronous Dynamic Programming

- The methods described so far used *synchronous* backups
 - ▷ All states are backed up at the same time
- *Asynchronous DP* backs up states individually, in any order
- Reduces computation time
- Guaranteed to converge if all states continue to be selected

2

Monte-Carlo Learning

Learning from **complete** episodes with
an **unknown** MDP



Monte-Carlo Learning

- **Model-free** – no need to know MDP transitions/rewards
- Learn directly from **complete** episodes of experience
- Simplest idea: use *empirical mean return* instead of *expected return*
 - ▷ $v_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$
 - ▷ $V(s) = S(s)/N(s)$
 - ▷ $V(s) \rightarrow v_{\pi}(s)$ as $N(s) \rightarrow \infty$
- Drawback: can only be applied to *episodic* MDPs - all episodes must terminate



Monte-Carlo Policy Evaluation

■ First-Visit MC Policy Evaluation

- ▷ Update estimate the **first time** a state s is visited in an episode.

■ Every-Visit MC Policy Evaluation

- ▷ Update estimate **every time** a state s is visited in an episode.

3

Temporal Difference Learning

Learning from **incomplete** episodes
with an **unknown** MDP



Temporal-Difference Learning

- **Model-free** – no need to know MDP transitions/rewards
- Learn directly from **incomplete** episodes of experience by **bootstrapping**
- Update a guess towards a guess



MC and TD

■ Goal: learn v_π online from experience under policy π

■ Incremental every-visit MC

▶ Update value toward **actual** return G_t

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

■ Simplest TD: TD(0)

▶ Update value toward **estimated** return $R_{t+1} + \gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha(\underbrace{R_{t+1} + \gamma V(S_{t+1})}_{\text{TD target}} - V(S_t))$$

TD target

TD error


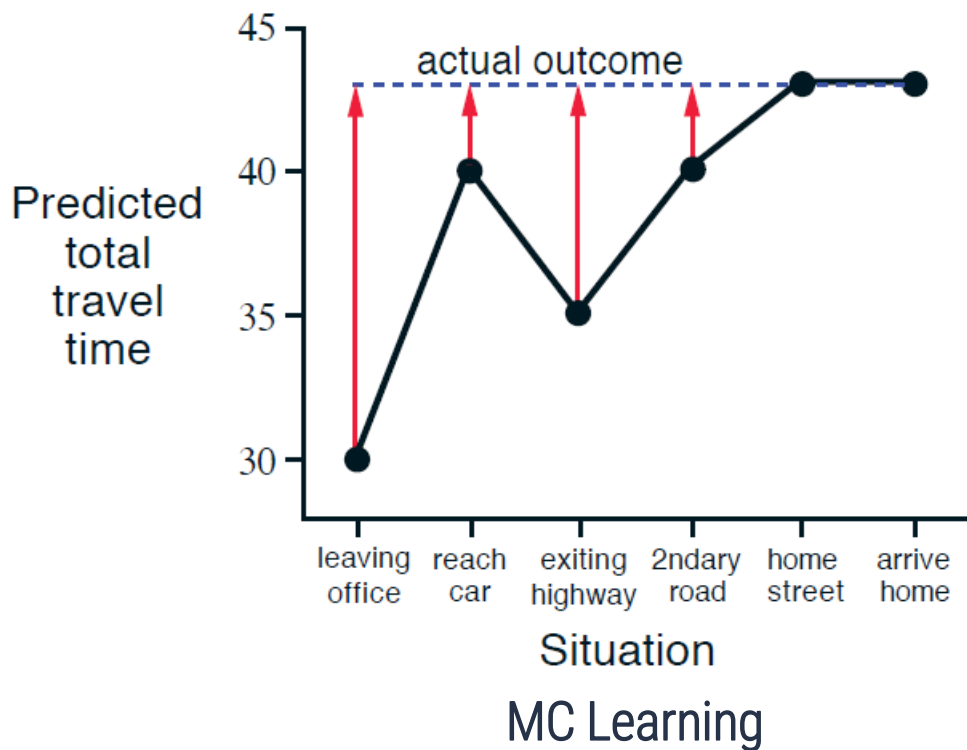


Driving Home Example

Situation/State	Elapsed Time (min)	Predicted Time Left (min)	Predicted Time: Total (min)
Leaving office	0	30	30
Reach car, raining	5	35	40
Exit highway	20	15	35
Behind truck	30	10	40
Home street	40	3	43
Arrive home	43	0	43



Driving Home Example





MC vs. TD

- TD can learn **before** knowing the final outcome
 - ▷ TD: can learn online after every step
 - ▷ MC: must wait until end of episode
- TD can learn **without** final outcome
 - ▷ TD: can learn from incomplete sequences
 - ▷ MC: can only learn from complete sequences
 - ▷ TD: works in non-terminating environments
 - ▷ MC: only works for terminating environments




MC vs. TD

- MC has **high variance, zero bias**
 - ▷ Good convergence properties
 - ▷ Not very sensitive to initial value
 - ▷ Very simple to use & understand
 - ▷ Does not exploit Markov property
- TD has **low variance, some bias**
 - ▷ More efficient than MC
 - ▷ TD(0) converges to $v_{\pi}(s)$
 - ▷ More sensitive to initial value
 - ▷ Exploits Markov property




Monte Carlo

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$



Temporal Difference


$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$






Dynamic Programming

$$V(S_t) \leftarrow \mathbb{E}_\pi [R_{t+1} + \gamma V(S_{t+1})]$$





Unified View





n -step Lookahead

Look n steps
into the
future

1-step TD
and TD(0)



2-step TD



3-step TD



...

n -step TD



...

∞ -step TD
and Monte Carlo





n -step Return

■ n -step return

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$


■ n -step TD learning

$$V(S_t) \leftarrow V(S_t) + \alpha \left(G_t^{(n)} - V(S_t) \right)$$



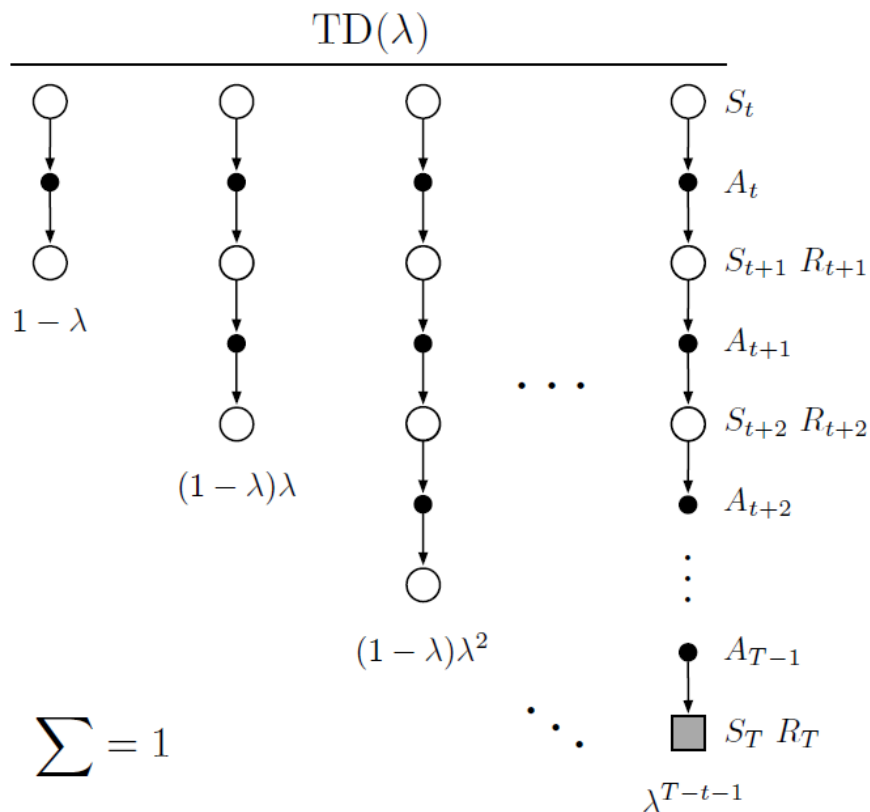
n -step TD

Average
RMS error
over 19 states
and first 10
episodes





TD(λ)



- Concept: get the best of all worlds
- Average n-step returns over all n, using weight $(1 - \lambda)\lambda^{n-1}$

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

- Forward-view TD(λ)

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t^\lambda - V(S_t))$$

4

Model Free Control

Optimizing with an unknown MDP



Model-free Reinforcement Learning

- Model-free prediction
 - ▷ **Estimate** the value function of an unknown MDP
- Model-free control
 - ▷ **Optimise** the value function of an unknown MDP



On and Off-Policy Learning

■ On-policy learning


- ▷ Learn about policy π from experience sampled from π

■ Off-policy learning

- ▷ Learn about policy π from experience sampled from μ



Generalized Policy Iteration (GPI)




Policy evaluation: Estimate v_π

Iterative policy evaluation

Policy improvement: Generate $\pi' \geq \pi$

Greedy policy improvement





Model-Free Policy Iteration

- Greedy policy improvement over $V(s)$ requires model of MDP


$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} \left(\mathcal{R}_s^a + \mathcal{P}_{ss'}^a V(s') \right)$$

- Greedy policy improvement over $Q(s, a)$ is model-free

$$\pi'(s) = \operatorname{argmax}_{a \in \mathcal{A}} Q(s, a)$$



Generalised Policy Iteration with Q




Policy evaluation: MC policy evaluation, $Q = q_\pi$

Policy improvement: Greedy policy improvement?



Example of Greedy Action Selection



JAKE-CLARK.TUMBLR

- There are two buttons:
- You press the **left** and get reward 0
 $V(\text{left}) = 0$
- You press the **right** and get reward +1
 $V(\text{right}) = +1$
- You press the **right** and get reward +3
 $V(\text{right}) = +2$
- You press the **right** and get reward +2
 $V(\text{right}) = +2$




ϵ -Greedy Exploration

- Simplest idea for ensuring **continual exploration**
- All m actions are tried with non-zero probability
- With probability $1 - \epsilon$, choose the greedy action
- With probability ϵ , choose an action at random

$$\pi(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon, & \text{if } a^* = \underset{a \in \mathcal{A}}{\operatorname{argmax}} Q(s, a) \\ \epsilon/m, & \text{otherwise} \end{cases}$$



Monte-Carlo Policy Iteration




Policy evaluation: MC policy evaluation, $Q = q_\pi$

Policy improvement: ϵ -greedy policy improvement

Note: we do this over multiple episodes



Monte-Carlo Control



Every episode (update with 'fresh' data):

Policy evaluation: MC policy evaluation, $Q \approx q_\pi$

Policy improvement: ϵ -greedy policy improvement




MC vs. TD Control

- TD has several advantages over MC:
 - ▷ Lower variance
 - ▷ Online
 - ▷ Incomplete sequences
- Use TD instead of MC in our control loop
 - ▷ Apply TD to $Q(\mathcal{S}, \mathcal{A})$
 - ▷ Use ϵ -greedy policy improvement
 - ▷ Update every time-step




Updating Q with SARSA



$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma Q(S', A') - Q(S, A))$$



On-Policy Control with Sarsa





Every **time-step**:

Policy evaluation: **Sarsa**, $Q \approx q_\pi$

Policy improvement: ϵ -greedy policy improvement



Windy Gridworld Example




standard
moves

- Reward = -1 per time step, until goal is reached
- Undiscounted




Sarsa on the Windy Gridworld






SARSA(λ)


Path taken




Action values increased by one-step Sarsa



Action values increased by 10-step Sarsa



Action values increased by Sarsa(λ) with $\lambda=0.9$





Q-Learning

- We now consider off-policy learning of action-values $Q(s, a)$
 - No importance sampling is required
 - Next action is chosen using behavior policy $A_{t+1} \sim \mu(\cdot | S_t)$
 - But we consider alternative successor action $A' \sim \pi(\cdot | S_t)$
 - And update $Q(S_t, A_t)$ towards value of alternative action
- $$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A') - Q(S_t, A_t))$$




Off-Policy Control with Q-Learning

- We now allow both behavior and target policies to **improve**
- The target policy π is **greedy** w/r $Q(s, a)$
- The behavior policy μ is **ϵ -greedy** w/r $Q(s, a)$
- Q-learning target simplifies as:


$$\begin{aligned} & R_{t+1} + \gamma Q(S_{t+1}, A') \\ &= R_{t+1} + \gamma Q\left(S_{t+1}, \underset{a'}{\operatorname{argmax}} Q(S_{t+1}, a')\right) \\ &= R_{t+1} + \max_{a'} \gamma Q(S_{t+1}, a') \end{aligned}$$



Q-Learning vs Sarsa



Sum of
rewards
during
episode



- Sarsa: on-policy, Q-Learning: off-policy
- We use Sarsa when we care about the agent's performance while learning (expensive robot)
- We use Q-learning when we don't mind the agent 'suffering'



Relationship Between DP and TD

	<i>Full Backup (DP)</i>	<i>Sample Backup (TD)</i>
Bellman Expectation Equation for $v_{\pi}(s)$	<p>A tree diagram representing a full backup. The root node is a circle labeled $v_{\pi}(s) \leftarrow s$. It has two children, both circles, representing next states. The left child is labeled $v_{\pi}(s') \leftarrow s'$. The root node is connected to the left child by an edge labeled a. The left child is connected to two leaf nodes (circles) by edges labeled r. The right child is also connected to two leaf nodes by edges labeled r.</p>	<p>A vertical sequence of three circles. The top circle is connected to the middle circle by an edge labeled r. The middle circle is connected to the bottom circle by an edge labeled a.</p>
Bellman Expectation Equation for $q_{\pi}(s, a)$	<p>A tree diagram representing a full backup for Q-values. The root node is a circle labeled $q_{\pi}(s, a) \leftarrow s, a$. It has two children, both circles, representing next states. The left child is labeled $q_{\pi}(s', a') \leftarrow s'$. The root node is connected to the left child by an edge labeled a. The left child is connected to two leaf nodes (circles) by edges labeled r. The right child is also connected to two leaf nodes by edges labeled r.</p>	<p>A vertical sequence of three circles. The top circle is labeled S, A. It is connected to the middle circle by an edge labeled R. The middle circle is labeled S'. The middle circle is connected to the bottom circle by an edge labeled A'.</p>
Bellman Optimality Equation for $q_{*}(s, a)$	<p>A tree diagram representing a full backup for Q-values. The root node is a circle labeled $q_{*}(s, a) \leftarrow s, a$. It has two children, both circles, representing next states. The left child is labeled $q_{*}(s', a') \leftarrow s'$. The root node is connected to the left child by an edge labeled a. The left child is connected to two leaf nodes (circles) by edges labeled r. The right child is also connected to two leaf nodes by edges labeled r.</p>	<p>A diagram representing a sample backup for Q-values. A central circle is connected to three leaf nodes (circles) by edges. The top leaf node is connected to the central circle by an edge labeled a. The other two leaf nodes are connected to the central circle by edges labeled r.</p>

5

Summary and Code

Examples of algorithms applied via
Python



Summary

- Dynamic Programming (DP)
 - ▷ Assumes perfect knowledge of the MDP
- Monte Carlo (MC) and Temporal Difference (TD) Learning
 - ▷ When the MDP is unknown
 - ▷ MC and TD both learn directly from experience
 - ▷ MC uses full episodes
 - ▷ TD uses bootstrapping
 - ▷ SARSA: On-Policy TD Control
 - ▷ Q-Learning: Off-Policy TD Control



Dynamic Programming Example

■ Gambler's Problem (from Sutton and Barto); Goal: reach exactly \$100



Step 1: BET



Step 2: TOSS COIN



Gain the
amount bet



Lose the
amount bet

Step 3: WIN/LOSE



Dynamic Programming Example

Value Iteration, for estimating $\pi \approx \pi_*$

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation
Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop:

```
|  $\Delta \leftarrow 0$   
| Loop for each  $s \in \mathcal{S}$ :  
|    $v \leftarrow V(s)$   
|    $V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$   
|    $\Delta \leftarrow \max(\Delta, |v - V(s)|)$   
until  $\Delta < \theta$ 
```

Output a deterministic policy, $\pi \approx \pi_*$, such that

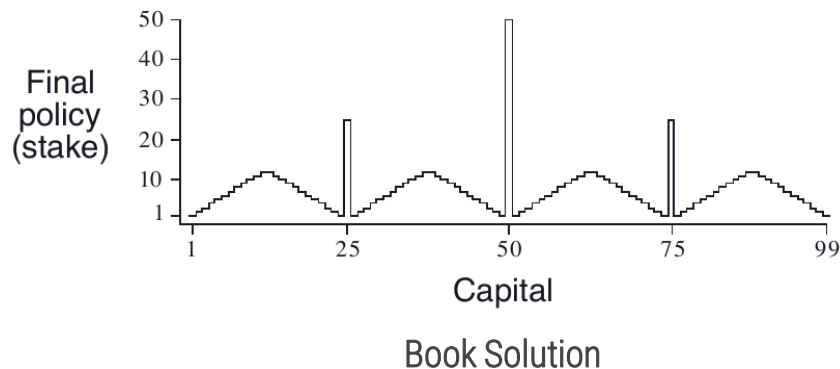
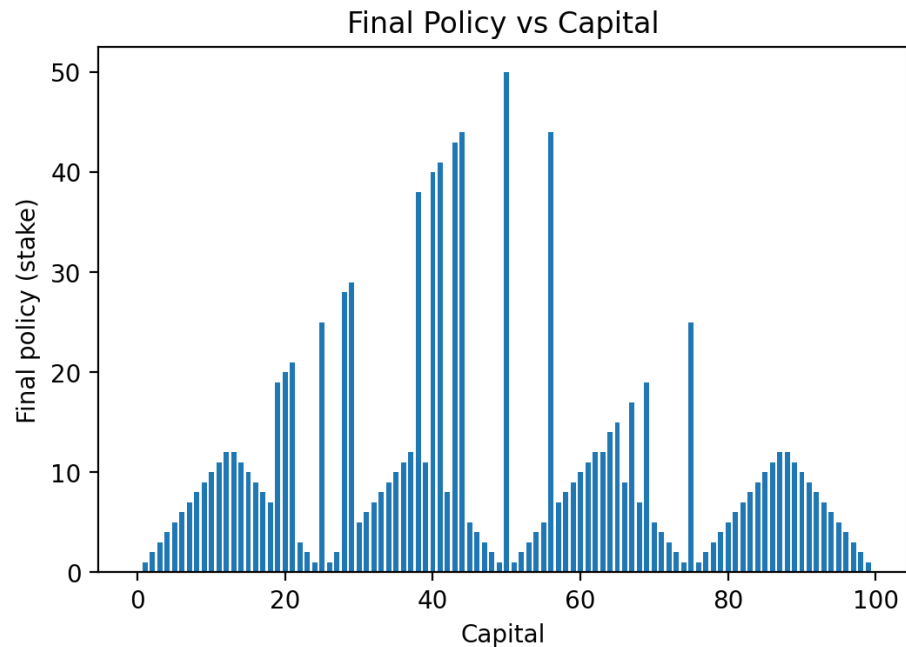
$$\pi(s) = \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$$

```
while True:  
    # Stopping condition  
    delta = 0  
    # Update each state...  
    for s in range(1, 100):  
        # Do a one-step lookahead to find the best action  
        A = one_step_lookahead(s, V, rewards)  
        best_action_value = np.max(A)  
        # Calculate delta across all states seen so far  
        delta = max(delta, np.abs(best_action_value - V[s]))  
        # Update the value function. Ref: Sutton book eq. 4.10.  
        V[s] = best_action_value  
    # Check if we can stop  
    if delta < theta:  
        break
```

```
# Create a deterministic policy using the optimal value function  
policy = np.zeros(100)  
for s in range(1, 100):  
    # One step Lookahead to find the best action for this state  
    A = one_step_lookahead(s, V, rewards)  
    best_action = np.argmax(A)  
    # Always take the best action  
    policy[s] = best_action
```




Dynamic Programming Example





Sneak-Peak: Actor-Critic for Mountain Car





Thank you!

- [UCL Course on Reinforcement Learning](#), David Silver
- [Reinforcement Learning – An Introduction](#), Sutton and Barto
- [Reinforcement Learning \(GitHub repo\)](#), Denny Britz

Special thanks to:

- Presentation template by [SlidesCarnival](#)
- Photographs by [Startup Stock Photos](#)